

VerusCite V1 Benchmark: Citation Verification Accuracy

Andrew Wheeler

2026-08-13

Table of contents

1	The Problem	1
2	Approach	2
2.1	Design Philosophy	2
3	Pipeline Architecture	2
3.1	Models Considered	3
3.2	Verification Categories	3
4	Benchmark	5
4.1	Extraction Results	5
5	Citation Checking Results	6
5.1	False Positives	7
5.2	Recall	8
5.3	Current Production Configuration	9
5.4	Cost Breakdown for Citation Checking	10
5.5	Population Estimates of Precision	11
6	Comparison between my tool and other approaches	12
7	Privacy	13
8	Limitations and Next Steps	13
9	Reproducibility	14
10	AI Use Disclosure	14

1 The Problem

Fabricated citations in academic papers have grown sharply since the widespread adoption of generative AI writing tools. A Lancet audit of 2.5 million biomedical papers found hallucinated reference rates rising 12-fold between 2023 and early 2026 (Topaz et al. 2026). Separate estimates project roughly 147,000 hallucinated citations across arXiv, bioRxiv, SSRN, and PMC in 2025 alone (Zhao et al. 2026).

Existing peer review and editorial safeguards largely fail to catch these errors before publication. The problem spans disciplines: biomedical research, AI/ML conferences, legal filings, and government reports.

While AI use is driving a surge in publications (see A. Wheeler 2026 for a practical overview), the same large language models that create this problem can also be used to verify citations at scale. VerusCite (veruscite-data.com) is a tool to automatically extract and verify citations from uploaded documents. This report documents the accuracy of the V1 verification pipeline against a hand-labeled ground truth corpus.

This tool is meant to be a quick and cheap approach to scan a bibliography for a paper, and determine if a citation is correct, has minor errors, or appears to have more serious issues, such as an AI generated hallucination. This paper serves as reference to the overall architecture and design of the project, as well as a source of false positive and negative rates on a benchmarked corpus of real papers.

2 Approach

2.1 Design Philosophy

Even in documents with hallucinated citations, the majority of references are legitimate. A tool with a high false positive rate – flagging real papers as hallucinations – would waste a reviewer’s time chasing false alarms rather than finding actual problems. VerusCite prioritizes **precision over recall**: it is better to miss some hallucinations than to incorrectly flag verified citations.

The target false positive rate for hallucination flags is below 1%. Recall for detecting actual hallucinations is around 70-80%. This trade-off means a human reviewer spends most of their time on citations that genuinely need attention rather than dismissing false reports.

The application is designed for human-in-the-loop review. Each citation’s verification status can be manually overridden, and the interface surfaces the reasoning behind each determination so reviewers can make informed decisions quickly.

Similar to work in Pangram, I believe that it is necessary to have the tool have as few of false positives as possible (Emi 2024). While this report will show that those false positives are not zero (any realistic tool should report a certain error rate), the design still prioritizes keeping them as low as possible.

3 Pipeline Architecture

VerusCite processes documents in two stages:

1. **Citation Extraction** – An LLM reads the document text and identifies individual references, parsing out structured fields (title, authors, year, journal, DOI).
2. **Citation Verification** – Each extracted citation is checked against external sources to determine whether it exists as described.

For citation extraction, the tool uses text searching across the document (if a PDF or word document is uploaded) to attempt to determine the location of the reference section. This tool currently works with reference sections combined in an article, it does not work with footnote style citations (as is common in law articles).

If uploading a bibliography file, such as a Bibtext `.bib` file, a `.ris` file, or an Endnote `.enw` file, the tool uses text tools to automatically extract out the reference information.

Once the reference section is identified, the reference section is then chunked into smaller portions, and uses a served LLM model (see later benchmark results) to extract out the citations into a structured sections. This means the LLM is only sent the references for the document, it is not sent other parts of the manuscript.

Each extracted citation is then checked independently. First the citation is checked against the Crossref database. If it is not found in Crossref, then an agentic tool uses web-search and fetch to verify whether the citation exists in external sources.

This independence means that in a document with some errors, other documents are not falsely assumed to be hallucinated. It also allows processing the documents in parallel, to return the results back faster.

3.1 Models Considered

Because the task involves web-search, when checking citations that are not verified via Crossref, this limits the potential pool of LLM providers. Here I consider two major companies:

- Gemini family models
- OpenAI family models

In particular, I focus on the cheaper models (flash-lite for Gemini and Nano/Luna for OpenAI) with minimal reasoning. I additionally consider these same models, but using the web search tools provided by *Perplexity*.

I do not consider Anthropic, as even their cheapest model (Haiku) is considerably more expensive than the models I evaluated here. I only partially evaluated AWS Nova-2 lite with their web search (latency was too long). Grok 4.3 did well in the benchmark tests but was considerably more expensive.

Extraction does not need web search, and so while I evaluated many other models on Bedrock (such as the open source GPT and Gemini models), they were not as accurate as the frontier served models. Extraction takes around 2 to 3 cents per document, so at this time I am not concerned about cost, although in the future likely even smaller open source models (especially if fine tuned) will be sufficient for the extraction task.

I would consider additional open source models in the future (such as DeepSeek flash) for citation checking, although it will be necessary to incorporate custom web search and fetch tools in conjunction with the served models. If the models are not provided out of the box by Perplexity.

3.2 Verification Categories

The checker assigns each citation one of four statuses:

Status	Meaning
Verified	Citation confirmed to exist via CrossRef or web search
Minor Error	Citation likely exists but has metadata discrepancies (wrong year, volume, pages)
Hallucination	No evidence the cited work exists as described
Not Found	Unable to confirm or deny – includes bare URL citations and paywalled/very recent works with insufficient search results

Many citations have minor errors (perhaps on the magnitude of 5% in my personal assessment). These include many innocuous things that are likely human introduced errors, like a year off or page numbers wrong.

The major different between a hallucination and a minor error is that hallucinations tend to have whole cloth differences that are difficult to explain. Such as journals entirely replaced. For example, Canessa et al. (2026) has the citation:

Elrod, Linda D. (2006). “A Move in the Right Direction? Best Interests of the Child Emerging as the Standard for Relocation Cases”. *Journal of the American Academy of Matrimonial Lawyers* 15, pp. 1-54.

This is a real paper, but should be cited as:

Elrod, Linda D. (2006). “A Move in the Right Direction? Best Interests of the Child Emerging as the Standard for Relocation Cases”. *Journal of Child Custody* 3, 29–61. https://doi.org/10.1300/J190v03n03_03

While the [Journal of the American Academy of Matrimonial Lawyers](#) does exist, there is no combination of volumes and page numbers that could reasonably be confused with the article in the *Journal of Child Custody*.

While no automated tool can never find 100% proof that a citation was hallucinated via a generative AI tool, it is difficult to construct a scenario where a bibliography had such errors manually generated, especially the volume and page numbers that do not exist in any comparable document for the Elrod citation.

Hallucinations can also include when multiple authors are incorrect, but author name errors are much more common, so they are less likely to be flagged. For an example of a minor error Mekonen (2026) has the reference:

33. Croft TMA, Allen CK, Arnold F, Assaf S, Balian S. *Guide to DHS Statistics: DHS-7 (version 2)*. Rockville, MD: ICF. 2020.

The actual author list includes an Aileen M.J. Marshall, whose initials were concatenated into the first author (see page 2 at https://dhsprogram.com/pubs/pdf/DHSG1/Guide_to_DHS_Statistics_DHS-7_v2.pdf).

Minor errors are incredibly prevalent. Any tool intended to do the same task needs to effectively distinguish between minor errors to be able to identify actual hallucinated citations.

4 Benchmark

The V1 validation corpus contains **2288** hand-labeled citations across **36 documents**. Sources are intentionally heterogeneous:

- Papers with known hallucinations identified by others on social media ([Chris Carothers on X](#), [David Buil-Gil on BlueSky](#))
- GPTZero’s NeurIPS 2025 audit (GPTZero 2026)

- Papers flagged by Reviewer3 (Reviewer3, n.d.)
- A ChatGPT deep-research output where all 12 citations are fabricated (Jacques et al. 2026)
- Clean papers (my own dissertation excerpts, open-access criminology, PLoS ONE articles)
- MDPI and preprint samples

The corpus spans categories in social sciences, physics, mathematics, as well as many different formats, including many different pre-print and journal format examples.

Of the 2288 citations, 1889 (82.6%) are verified, 204 (8.9%) have minor errors, 141 (6.2%) are hallucinations, and 54 (2.4%) are not found. The corpus skews heavily toward verified citations (as real documents do), which makes false positive rate the critical metric.

There are two different metrics a user should be interested in. First is whether the tool can actually extract out the references accurately. This is non-trivial in and of itself – a skills based approach submitting a large document with an extensive bibliography is likely to generate more errors in this step.

The second metric is in terms of false positives (e.g. a good citation is flagged as a hallucination), and recall rates (of all the hallucinations, how many does the tool capture). Below are those metrics for this corpus on the current VerusCite tool across different LLM providers.

4.1 Extraction Results

Citation extraction uses an LLM to parse the document text and identify individual references. The table below shows extraction accuracy against the ground truth. There are two types of errors that can occur, you can miss a reference, or the tool can add in a reference.

Table 1: Extraction accuracy by model

Model	Real	Correct	Missing	Extra	Rate (%)	Wall min	Docs/min	Cost
gemini-3.1-flash-lite	2288	2285	3	3	99.9	4.9	7.4	1.15
gpt-5.4-nano	2288	2285	3	4	99.9	12.2	3.0	0.88
gpt-5.6-luna	2288	2284	4	6	99.8	8.2	4.4	0.84

Extraction is very accurate across models. Gemini 3.1 Flash Lite (using pypdfium2 to read in the text, we additionally considered litemap with markdown) correctly extracts over 99.8% of citations at under \$0.03 per paper for the full corpus. The few “extra” rows are typically page footers, footnotes, or duplicate listings that the LLM picks up in addition to the actual bibliography.

I also tested `litemap` with the markdown export format to extract out the text from PDFs, which tended to produce slightly more artifacts – it does not distinguish line breaks for individual

citations as cleanly as pypdfium2, producing a few more missed or merged entries. Results were largely similar overall (and I may migrate to using liteparse with markdown in the future).

Both the extraction part of the pipeline and the citation check part of the pipeline have fallback models, as it is common for these LLM providers to have downtime with models. The current production version of VerusCite uses 3.1 flash lite (served by Google) as the primary extraction model (due to both accuracy and cost). The fallback model is currently gpt-5.6-luna. While OpenAI (both the 5.4 nano and 5.6 luna models) are somewhat cheaper, due to the lower accuracy and time, it is not used at this point. (The cost only saves around 1 cent per paper.)

Generally the prompts are constructed in a way that advanced reasoning is not necessary. So going up to larger reasoning models (and expanding the reasoning budget) does not result in higher accuracy. For OpenAI, you can see that the recent luna model did not result in any more accurate results than nano.

5 Citation Checking Results

The citation checking results break down the benchmark against both the model, and the provider. When Perplexity is the provider, it means the tool called Perplexity servers (using their internal web search), whereas if the provider is gemini it uses Googles web search tool, or OpenAI uses its internal web search tool. Otherwise all prompts are the same across each of the models.

The citation checking results input in correct data – extraction and checking are done in independent benchmarks. So each of the 2288 citations, it is independently run though the checking benchmark, swapping out the model used.

Where the same model/provider appears more than once, those are **independent full-corpus re-runs** (different dates), not duplicates of one output. Repeat runs are intentional: agentic web search and LLM classification are stochastic, so multiple runs show run-to-run variance in false positives, recall, cost, and wall time. Accuracy tables include a **Date** column so repeats are distinguishable.

Table 2: Citation Checking accuracy by model/provider (repeat rows are independent re-runs)

Model	Provider	Date	FP Hall.	FP		Hall. Recall	NV Recall
				Minor			
gemini-3.5-flash-lite	gemini	2026-07-29	0.1% (2)	0.8% (15)		63.1% (89/141)	71.9% (286/398)
google/gemini-3.1-flash-lite	perplexity	2026-07-30	0.4% (7)	0.7% (13)		69.5% (98/141)	82.7% (329/398)
google/gemini-3.1-flash-lite	perplexity	2026-08-03	0.4% (7)	1.2% (22)		64.5% (91/141)	80.2% (319/398)
google/gemini-3.1-flash-lite	perplexity	2026-08-13	0.4% (7)	1.1% (20)		47.5% (67/141)	80.2% (319/398)

Table 2: Citation Checking accuracy by model/provider (repeat rows are independent re-runs)

Model	Provider	Date	FP Hall.	FP	Hall. Recall	NV Recall
				Minor		
google/gemini-3.1-flash-lite	perplex-ity	2026-08-13	0.6% (11)	0.9% (17)	68.8% (97/141)	79.1% (315/398)
gpt-5.4-nano	openai	2026-07-30	0.2% (3)	2.2% (42)	60.3% (85/141)	88.4% (352/398)
gpt-5.6-luna	openai	2026-07-30	0.2% (3)	2.5% (47)	74.5% (105/141)	88.7% (353/398)
gpt-5.6-luna	openai	2026-08-03	0.1% (2)	2.8% (52)	73.8% (104/141)	89.2% (355/398)
gpt-5.6-luna	openai	2026-08-13	0.2% (4)	2.4% (45)	72.3% (102/141)	88.7% (353/398)
openai/gpt-5.4-nano	perplex-ity	2026-07-30	0.4% (7)	2.5% (47)	33.3% (47/141)	84.9% (338/398)
openai/gpt-5.6-luna	perplex-ity	2026-08-01	0.2% (4)	2.2% (41)	44.7% (63/141)	87.7% (349/398)
openai/gpt-5.6-luna	perplex-ity	2026-08-13	0.2% (3)	2.0% (38)	44.0% (62/141)	87.7% (349/398)
perplexity/nemotron-3.5-lightning-30b-a3b	perplex-ity	2026-08-13	0.1% (1)	0.5% (10)	30.5% (43/141)	63.6% (253/398)

5.1 False Positives

For key metrics, there are a total of 1889 correct citations in the corpus. So false positives (FP) rates are of those 1889 correct citations, how many were falsely flagged by the tool.

The current FP rates for hallucinations are well under 1%, and vary between 1 to 11 total false positives across the different models in this corpus (and can shift by a few counts between re-runs of the same model). Hallucination false positives are often attributable to idiosyncratic web search results. (The agentic tools often will return a response, even if the web search tool is currently inaccessible.) Verified false-positive rates on Perplexity are broadly similar to the same model on the provider’s own web-search stack, but **hallucination recall is not**: Perplexity-hosted OpenAI models (`openai/gpt-5.4-nano`, `openai/gpt-5.6-luna`) under-detect true hallucinations relative to OpenAI direct (see Recall below).

For an example of a false positive hallucination in this particular run, my dissertation has the citation:

Pearl, J. (2000). Causality: models, reasoning and inference, Volume 29. Cambridge Univ Press.

And 5.4 nano lists as reasoning for classifying this as a hallucination:

Upstream verdict: book exists and matches author/title/year/publisher, but the cited “Volume 29” element is not supported by the primary publisher/library records

surfaced in the tool results and is treated as fabricated metadata; therefore status mapped to hallucination.

This behavior is not consistent, either within 5.4. nano or across other models. (This most often returns a “minor error”, although my ground truth I have this as “verified”. Some sources list the published year as 2001, causing additional issues beyond just the “Volume 29” addition.) I have intentionally avoided K-shot examples in the prompt, so there is as little leakage as possible and the prompts should better generalize to out of sample data. But given the stochastic nature of LLMs, some errors will ultimately occur.

FP rates for minor errors are more prevalent, being close to 1% for the Gemini models, but over 2% for the OpenAI models. Production defaults are summarized in [Current Production Configuration](#): primary checker is **Perplexity-hosted Gemini 3.1 Flash Lite**, with **OpenAI direct gpt-5.6-luna** as the fallback.

The general approach I took was to evaluate when both OpenAI and Google models returned false positives. While these do happen in the corpus, in some cases they are inevitable, as CrossRef or other online data provides conflicting information. For one example citation:

[CS22] T. Cieřla and M. Sabok. Measurable Hall’s theorem for actions of abelian groups, J. Eur. Math. Soc., 24 (2022), 2751-2773 (cit. on pp. 3, 44)

Consistently produces a “minor error” category in the tool due to the citation year. Going to the website, the citation year is correctly 2022. But [crossref lists the publication year as 2021](#).

Many of the false positives are minor errors like this (and likely some of the labels in the ground truth should be updated – feel free to contact me if you believe an articles classification is not correct in the ground truth). False positives for not found tend to be due to unreliable web search, and so are intermittent and not consistent due to the tool.

5.2 Recall

Recall is the proportion of the true errors that are captured by the current tool. These are more variable across the different tools, with OpenAI models having greater recall (which comes with more false positives). Hallucination recall for the Gemini configurations is typically in the **60–70%** range, but one Perplexity 3.1 flash-lite re-run (2026-08-13) dropped to **47.5%** (67/141) while a later same-day re-run recovered to **68.8%** (97/141). Across three independent full-corpus runs, **OpenAI direct gpt-5.6-luna** lands at about **72–75%** (102–105/141), while production **Perplexity google/gemini-3.1-flash-lite** is usually about **65–70%** (91–98/141) aside from that one low re-run. The same luna weights on **Perplexity (openai/gpt-5.6-luna)** stay near **44–45%** hallucination recall (62–63/141) across the 2026-08-01 and 2026-08-13 runs, with not-verified recall still high (~88%) because many misses land in minor error or not found rather than verified. Perplexity **openai/gpt-5.4-nano** shows the same pattern even more sharply

(33% hallucination recall). A Perplexity **Nemotron 3.5 Lightning** run is much cheaper but weaker on errors (**30.5%** hallucination recall, **36.9%** minor-error recall). Provider web-search quality matters as much as the base model for this task.

That re-run spread is the main reason tables keep multiple rows for the same model: FP minor for Perplexity 3.1 flash-lite moved from 0.7 (13) to 1.2% (22) between the July and early-August runs, then sat at 0.9–1.1% on the 2026-08-13 pair, while OpenAI luna hallucination FP stayed at 0.1–0.2% and hallucination recall stayed near 72–75%. Cost and wall time also drift with provider load and search volume (see the cost table).

Minor error recall is similarly lower for Gemini direct 3.5 flash-lite, at 52.7%. OpenAI direct has higher recall (often higher than 70%) for minor errors. Not found recall is near perfect across all models.

The final category, not verified, collapses the categories of minor error, hallucination, and not found. So if many hallucinations were classified into minor error, the direct recall rates would be low, but not verified (which will typically trigger a human review) would still be high. These are consistently over 80% across Gemini and OpenAI runs, with the exception of Gemini 3.5 flash lite (mostly low minor-error recall) and Nemotron (63.6% not-verified recall).

5.3 Current Production Configuration

Production VerusCite uses the following defaults for citation checking:

Role	Provider	Model
Primary checker	Perplexity	<code>google/gemini-3.1-flash-lite</code>
Fallback checker	OpenAI (direct)	<code>gpt-5.6-luna</code>

Why Perplexity 3.1 Flash Lite as primary. Google’s Gemini web-search stack caps searches at **1,500 per day across all tiers**, which is too low for multi-user production load (a single large paper can consume dozens of searches after Crossref misses). Perplexity does not impose that daily search cap, so agentic verification can keep running under concurrent documents. On this corpus, Perplexity `google/gemini-3.1-flash-lite` re-runs keep hallucination FP at or under 0.6% and not-verified recall around **79–83%**.

Gemini **direct** 3.5 flash-lite remains the fastest and cheapest full-corpus configuration in the table below (~26s wall-equivalent per paper under concurrent batching, ~\$0.41/paper), but the search quota makes it unsuitable as the default production path. It is retained in the benchmark for comparison.

Why OpenAI luna as backup. Direct OpenAI `gpt-5.6-luna` is the fallback when Perplexity is down or returns persistent errors. Across three re-runs it has the highest hallucination recall among configurations tested (~**72–75%**). Later re-runs are cheaper than the July/early-August pair (the 2026-08-13 OpenAI luna run was about **\$15 / \$0.42 per paper**). The

Perplexity-hosted OpenAI variants (`openai/gpt-5.4-nano`, `openai/gpt-5.6-luna`) are **not** used as production fallback: they under-detect true hallucinations relative to OpenAI direct (see Recall above). Nemotron on Perplexity is also not used in production: it is the cheapest full-corpus row (~\$6.53) but misses most real errors.

Extraction production remains Gemini 3.1 Flash Lite (Google) with OpenAI fallback, as described in the Extraction Results section.

5.4 Cost Breakdown for Citation Checking

For the 2288 total citations, approximately 1310 citations were verified via Crossref after extraction (about 57%) for each of the models. These are largely automated, and so incur no additional LLM cost. When those fail however, an agent based LLM tool needs to use web search and fetch to identify whether the citation is correct. Thus costs incur for both token usage as well as web search. Web search costs \$5 per 1000 searches on Perplexity, \$7 per 1000 searches for OpenAI, and \$14 per 1000 searches for Google.

The prompts are generally short enough that token caching does not occur at all for the Gemini models (needs over 4000 tokens). Some token caching does occur for OpenAI (although these costs do not include that), but it is relatively small (and the majority of token costs are output).

Table 3: Cost and runtime breakdown per full-corpus checker run (36 papers; repeat rows are independent re-runs)

Model	Provider	Date	Token Cost (USD)	Search Cost (USD)	Total (USD)	Per Paper (USD)	Sec/paper	Wall min
gemini-3.5-flash-lite	gemini	2026-07-29	6.14	8.44	14.59	0.41	128.9	15.7
google/gemini-3.1-flash-lite	perplexity	2026-07-30	12.15	7.72	19.87	0.55	406.2	51.0
google/gemini-3.1-flash-lite	perplexity	2026-08-03	13.43	8.94	22.36	0.62	448.7	57.9
google/gemini-3.1-flash-lite	perplexity	2026-08-13	7.85	4.2	12.06	0.34	339.6	41.9
google/gemini-3.1-flash-lite	perplexity	2026-08-13	6.42	4.21	10.63	0.3	185.9	22.7
gpt-5.4-nano	openai	2026-07-30	7.17	16.28	23.45	0.65	445.5	54.4
gpt-5.6-luna	openai	2026-07-30	6.84	11.49	18.33	0.51	328.6	39.9
gpt-5.6-luna	openai	2026-08-03	6.9	11.53	18.43	0.51	386.2	47.1
gpt-5.6-luna	openai	2026-08-13	3.26	11.87	15.13	0.42	263.4	32.2
openai/gpt-5.4-nano	perplexity	2026-07-30	7.34	4.84	12.18	0.34	243.3	29.7

Table 3: Cost and runtime breakdown per full-corpus checker run (36 papers; repeat rows are independent re-runs)

Model	Provider	Date	Token Cost (USD)	Search Cost (USD)	Total (USD)	Per Paper (USD)	Sec/pa- per	Wall min
openai/gpt-5.6-luna	per- plexity	2026- 08-01	11.21	8.76	19.98	0.56	291.6	35.7
openai/gpt-5.6-luna	per- plexity	2026- 08-13	4.04	6.04	10.07	0.28	319.4	38.9
perplexity/nemotron-3.5- lightning-30b-a3b	per- plexity	2026- 08-13	1.66	4.88	6.53	0.18	730.2	88.7

Sec/paper is the mean document wall time from each run’s `document_metrics.csv` (how long one paper takes end-to-end under the batch concurrency settings). **Wall min** is full-corpus elapsed time with up to five documents in parallel. Repeat rows again show variance: Perplexity 3.1 flash-lite was about **\$19.87 / 51 wall min** and **\$22.36 / 58 wall min** in July/early August, then **\$12.06 / 42 min** and **\$10.63 / 23 min** on 2026-08-13; OpenAI direct luna was about **\$18.3–\$18.4** with wall time **40–47 min**, then **\$15.13 / 32 min** on 2026-08-13.

The newer Gemini-direct 3.5 flash lite model has lower token costs and the shortest per-paper times, but is not the production primary for the web-search quota reasons above. Across this corpus, costs are typically around 30–60 cents per paper, with often more than half of the cost devoted to web search. Production primary (Perplexity 3.1 flash lite) on the latest re-run is about **\$0.30/paper** and roughly **23 minutes** full-corpus wall (five documents in parallel). OpenAI direct luna fallback on the latest re-run is about **\$0.42/paper** and **32 minutes** full-corpus wall. Perplexity `openai/gpt-5.6-luna` also got cheaper on 2026-08-13 (~\$10.07 full corpus) while keeping the same weak hallucination-recall pattern. Perplexity Nemotron is the lowest total cost (~\$6.53) but is not accurate enough for production.

5.5 Population Estimates of Precision

The ground truth sample I have collected is likely not representative of the overall population of scholarly papers. I have intentionally selected examples papers that have high rates of hallucinated citations. In Topaz et al. (2026) and Zhao et al. (2026), they estimate that pre-print servers and PubMed papers currently have rates of hallucinated citations ranging from 0.2% to almost 1.9%.

So for an hypothetical example, if false positive rates for hallucinations in my tool are 0.5%, and the population prevalence rate of hallucinations are 1%, and my recall rates are 60%, what are the estimates for precision in the population? The estimated precision in the population will be only slightly over 50%. So out of 100 citations, 1 is a real hallucination, and 99 are not. Overall I capture that real hallucination 60% of the time, so 0.6 true positives. Of the 99 I have 0.99×0.005 false positives, so approximately 0.5 false positives. And then the precision is then $0.6 / (0.5 + 0.6)$, which is slightly over 50% (Gigerenzer 2011).

If true rates of hallucinations are lower, my population precision will also be lower. Increasing recall could help, but even if recall is increased to 80% precision is still only 60%.

As such, it will be necessary for humans to carefully review the output of the tool. If an article with a single hallucinated is identified, it does not for sure indicate that AI was used in the production of the bibliography. If multiple references are flagged as hallucinations though, it is more likely there are substantive issues with the work that a dutiful editor should investigate.

Ultimately they overall rate of 1% hallucinations are a mixture of papers – most scientists likely do a good job, but there are likely a few bad apples that use AI to generate whole slews of papers. And those individuals may have hallucination rates in over 20% of their citations (Walters and Wilder 2023). Although with the improvement of the generative AI tools, that will likely decrease over time.

6 Comparison between my tool and other approaches

For some differences between my work and Topaz et al. (2026) and Zhao et al. (2026), my tool can flag as hallucinations when authors or journals are clearly wrong. Both Topaz et al. (2026) and Zhao et al. (2026) focus on title matches (and mismatches) only. The prior example I gave in Canessa et al. (2026) would not be classified as a hallucination in either of these tools, as the title exists – it just clearly swapped out an incorrect journal.

Additionally both Topaz et al. (2026) and Zhao et al. (2026) use google scholar as a secondary reference source. I do not consider this as valid, as I have personally seen google scholar reference hallucinated citations, see [this example of a hallucinated reference](#), the citation

[5] J. Smith, Spatial analysis of crime patterns using GIS-based decision support tools, J. Quant. Criminol. 38 (2) (2022) 125–148

that is in Kebede et al. (2026) in my ground truth, and actually links to the (false) google scholar citation.

These methods focus on title matches. So for example, Topaz et al. (2026) give the example as a *not* hallucinated title:

For example, a reference listed as *Depression and anxiety in young adults with ID* corresponds to the real indexed title *Depression and anxiety symptoms during the transition to early adulthood for people with intellectual disabilities* and is probably a reference error, not a fabrication.

My current system, while calling stochastic LLMs, reliably classifies this as a *minor error* (across all model permutations).

Other current tools, like GPTZero or Reviewer3, do not publish how they exactly approach citation checking. While Reviewer3 lists a [reference set of metrics](#), they do not share the actual

citations. (It appears these are simulated, as the Walters and Wilder (2023) paper they cite has 636 citations, and Reviewer3’s reference set only has 476 citations.)

The articles additionally focus on peer reviewed literature, and do not evaluate hallucinations in non journal articles (grey literature). My benchmark includes the full corpus in the reference papers, including blog posts, working papers, datasets, technical reports, law cases, etc. I have included the not found category for the scenario where a citation cannot be reasonably identified from web based resources, but including these sources will likely increase false positive rates, as how they are even supposed to be correctly cited can be difficult to know exactly.

7 Privacy

VerusCite uses only zero data retention (ZDR) models from all providers. Uploaded documents are not used for model training by any third party.

Additionally, the pipeline minimizes what is sent to external LLMs. Citation extraction uses text search on the locally-extracted PDF text first – only the citation strings themselves are sent to an LLM for structured parsing. During verification, only the parsed citation metadata (title, authors, year) is sent to web search, not the full document text. The full PDF content never leaves the server.

8 Limitations and Next Steps

The V1 corpus is heterogeneous by design, but 36 documents is a limited sample. The next evaluation (V2) will use the same prompts and pipeline but a **completely different validation set** – different documents, different domains, different sources of known hallucinations. This ensures the system is not overtrained on the particular examples in V1.

Known limitations of the current evaluation:

- Ground truth labels for some documents have ambiguous categories that likely could be reasonably changed. This is particularly true for the minor error category.
- “Not found” is an inherently ambiguous category – some citations are bare URLs, others exist but are difficult to locate via web search (paywalled, very recent, or in non-English databases).
- The corpus over-represents arXiv preprints relative to other domains (law, humanities, clinical medicine).

9 Reproducibility

All data for this report is public at github.com/apwheelee/veruscite-data. The `v1/` directory contains:

- `ground_truth.csv` – hand-labeled citations with `expected_status`
- `extraction_run/` – raw extraction outputs per model
- `checking_run/` – raw checker outputs per model
- `manifest.json` – run IDs used in this report

To regenerate this report:

```
pip install -r requirements.txt
quarto render report.qmd
```

10 AI Use Disclosure

This paper was prepared with AI assistance. Drafting and earlier iterations used **Claude Opus 4.6** (Anthropic), reviewing prior works by Andrew Wheeler (see A. P. Wheeler (2026) for that workflow). Additional edits, including the Current Production Configuration section, multi-run variance notes for independent checker re-runs (Perplexity `google/gemini-3.1-flash-lite` and OpenAI `gpt-5.6-luna`), per-paper runtime in the cost table, and the 2026-08-13 re-runs plus Nemotron row, were assisted by **Grok 4.5 / 4.6** (xAI). Ground-truth labels and final review are done by myself (Andrew P. Wheeler). All errors are my own.

References

- Canessa, Stella, Gordon B Dahl, Anna Hasselqvist, et al. 2026. *Life After Divorce: Effects of Joint Custody on Parents and Children*. Working Paper No. 35482. Working Paper Series. National Bureau of Economic Research. <https://doi.org/10.3386/w35482>.
- Emi, Bradley. 2024. *All about False Positives in AI Detectors*. <https://www.pangram.com/blog/all-about-false-positives-in-ai-detectors>.
- Gigerenzer, Gerd. 2011. “What Are Natural Frequencies?” *BMJ* 343.
- GPTZero. 2026. *GPTZero Finds 100 New Hallucinations in NeurIPS 2025 Accepted Papers*. <https://gptzero.me/news/neurips/>.

- Jacques, Scott, Andrew Wheeler, and Joshua Gerstenfeld. 2026. “Open Access, Generative Artificial Intelligence, and the Criminology Evidence Base.” *Evidence Base* 1 (2): 2658591. <https://doi.org/10.1080/30679125.2026.2658591>.
- Kebede, Hailu, Mohammed Motuma Assen, and Merid Abadi Sharew. 2026. “Crime Hotspot Analysis and Mapping Using Geospatial Technology in Dessie City, Ethiopia.” *Next Research* 5: 101303. <https://doi.org/10.1016/j.nexres.2025.101303>.
- Mekonen, Enyew Getaneh. 2026. “Prevalence and Associated Factors of Intimate Partner Violence Against Reproductive-Age Women in Africa and Asia Regions: Insights from 2022–2024 DHS Datasets.” *PLOS ONE* 21 (7): 1–16. <https://doi.org/10.1371/journal.pone.0353669>.
- Reviewer3. n.d. *Reviewer3: Live arXiv Reference Checking*. <https://reviewer3.com/live/arxiv>.
- Topaz, Maxim, Nir Roguin, Pallavi Gupta, Zhihong Zhang, and Laura-Maria Peltonen. 2026. “Fabricated Citations: An Audit Across 2.5 Million Biomedical Papers.” *The Lancet* 407 (10541): 1779–81. [https://doi.org/10.1016/S0140-6736\(26\)00603-3](https://doi.org/10.1016/S0140-6736(26)00603-3).
- Walters, William H., and Esther Isabelle Wilder. 2023. “Fabrication and Errors in the Bibliographic Citations Generated by ChatGPT.” *Scientific Reports* 13 (1): 14045.
- Wheeler, Andrew. 2026. *Large Language Models for Mortals: A Practical Guide for Analysts with Python*. Crime De-Coder.
- Wheeler, Andrew P. 2026. *Using Claude Code to Help Me Write*. <https://andrewpwheeler.com/2026/03/20/using-claude-code-to-help-me-write/>.
- Zhao, Zhenyue, Yihe Wang, Toby Stuart, Mathijs De Vaan, Paul Ginsparg, and Yian Yin. 2026. *LLM Hallucinations in the Wild: Large-Scale Evidence from Non-Existent Citations*. <https://arxiv.org/abs/2605.07723>.